Introduction to Time Series Analysis

OC Data Driven Insights

Eric Weber and Sarah Nooravi

August 17, 2018

Resources

- ▶ Forecasting: Principles and Practice by Rob J. Hyndman and George Athanasopoulos
- ▶ Time Series Analysis and Its Applications with R Examples by Robert H. Shumway and David S. Stoffer

Outline

- ▶ What is a time series? Time series model?
  - ▶ Time Series Examples
- ▶ Fundamental building blocks of time series
  - ▶ White noise
  - ▶ Moving average
  - ▶ Random walk
- ▶ Second order properties and stationarity
- ▶ When/why linear regression fails

What is a Time Series?

A sequence of values of a variable at **equally spaced** intervals with a **natural temporal ordering**.

```
> require(astsa)
> window(jj, 1960, c(1965,4))
```

```
##      Qtr1 Qtr2 Qtr3 Qtr4
## 1960 0.71 0.63 0.85 0.44
## 1961 0.61 0.69 0.92 0.55
## 1962 0.72 0.77 0.92 0.60
## 1963 0.83 0.80 1.00 0.77
## 1964 0.92 1.00 1.24 1.00
## 1965 1.16 1.30 1.45 1.25
```

What is a Time Series?

```
> plot(jj, type="o", ylab="Quarterly Earnings per Share")
```

What is a Time Series?

```
plot(gtemp, type="o", ylab="Global Temperature Deviations")
```

Key Assumption

Can't assume consecutive observations are independent

```
> jj_v <- unclass(jj)
> lag <- lag(jj_v, 1)
> lm <- lm(jj_v ~ lag)
> cor.test( ~ jj_v + lag,
+           method = "pearson",
+           conf.level = 0.95)


##
##   Pearson's product-moment correlation
##
## data:  jj_v and lag
## t = 25.987, df = 81, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to
## 95 percent confidence interval:
##  0.9159239 0.9641248
## sample estimates:
```

Key Assumption

Can't assume consecutive observations are independent

```
> plot(jj_v, lag, ylab="lag(y)", xlab="y")
```

Why Care about Time Series?

- ▶ They are everywhere!
- ▶ Economics (stock market, unemployment, etc)
- ▶ Social sciences (population series like birth rates or school enrollments)
- ▶ Epidemiology (influenza outbreaks)
- ▶ Medicine (blood pressure measurements)

Working with Time Series

- ▶ Time series analysis
  - ▶ Analyzing observed data
  - ▶ Focus on characteristics of the data
  - ▶ Explanatory focus
- ▶ Time series forecasting
  - ▶ Generating a model
  - ▶ Predictive focus

Some Additional Terminology

- ▶ Stochastic process
  - ▶ A sequence of random variables
  - ▶ Example: flipping a coin
- ▶ Sample path
  - ▶ Sample path of a stochastic process
  - ▶ One sample from a stochastic process
  - ▶ For example: HTHHTT (six coin flips)
- ▶ A stochastic process can generate MANY sample paths (infinitely many)

Fundamental Stochastic Processes

- ▶ White noise
- ▶ Moving average
- ▶ Random walks

White Noise

- Fundamental building block of other stochastic processes
- Key: no correlation between observations
- Y_t is called white noise (process looks like white light of spectrometers)

White Noise: Definition

- Assume $e_t \sim N(0, 1)$ is a collection of IID random variables all following a normal distribution
- Normal distribution? Anyone? Anyone at all?
- Define $Y_t = e_t$ for all t
- A white noise process satisfies:
  - Mean(Y_t) = constant
  - Var(Y_t) = constant
  - Auto-covariance = 0

▶ Since the data is random (no pattern) there is no point in modeling

# White Noise: 1 Sample Path

White Noise: Many Sample Paths

Moving Average

- Big difference: there IS correlation between observations
- Assume $e_t \sim N(0, 1)$ is a collection of IID random variables all following a normal distribution
- Define $Y_t = \frac{e_t + e_{t+1}}{2}$ for all t
- Y_t is a type of moving average stochastic process

# Moving Average: 1 Sample Path

Moving Average: Many Sample Paths

Random Walk

- Assume $e_t \sim N(0, 1)$ is a collection of IID random variables all following a normal distribution
- Define $Y_t = e_1 + e_2 + ... + e_t$ for all t
- Y_t is a random walk

Random Walk: 1 Sample Path

Random Walk: Many Sample Paths

Why Care about Stochastic Processes?

- We have ONE sample path of observations
- From this sample path we intend to infer the stochastic process that generated it
- We are NOT fitting lines to the data
- We are understanding the sample paths the process could take

- We observe a process to a specific point
- We determine what sample paths are likely as we look later in time

Stationarity

- ▶ Heavily mathematical definition
- ▶ Essence of it:
    - ▶ The mean (or expected value of the stochastic process) is constant
    - ▶ Covariance (or corelation between points only depends on distance between points and not on the value of t)
- ▶ Stationarity is an extremely common assumption in time series modeling!

Time Series: Linear Trend

- Assume stochastic process is $Y_t = \beta_0 + \beta_1 t + X_t$
- Assume $E[X_t] = 0$ and X_t is stationary
- This is just like linear regression so let's use OLS

Time Series: Linear Trend

Time Series

▶ Linear regression recovers the true parameters (close)

```
## # A tibble: 2 x 5
##   term         estimate std.error statistic  p.value
##   <chr>           <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)      30.7     11.3       2.73 7.57e- 3
## 2 t                 3.13     0.194     16.2 2.13e-29
```

- ▶ What if we fit a random walk Y_t with a linear model?
- ▶ Put another way: what if we regress a random walk on time with a linear model?
- ▶ Good choice, bad choice?

Linear Model for a Random Walk

Linear Model for a Random Walk

▶ Linear regression fails. Hard
▶ Significant trend when there is NOT one
▶ Reminder: there is not a trend because the expected value of Y_t is 0

```
## # A tibble: 2 x 5
##   term        estimate std.error statistic  p.value
##   <chr>          <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)   -150.      22.9     -6.53 2.97e- 9
## 2 t               13.1      0.394    33.2  4.34e-55
```

What Happened?

- ▶ Spurious regression
- ▶ DO NOT regress non-stationary time series
- ▶ Our mistake
  - ▶ Coefficient is okay
  - ▶ Standard error is underestimated
- ▶ A random walk does not follow required linear regression assumptions!
  - ▶ Which key assumption does it not follow?